Case Study: MakeMyTrip AI Trip Assistant

Role: Product Manager (AI Product Case Study) · Type: AI-native feature design, no-code prototype, pricing & GTM strategy Links: Airtable Base *(replace with your real Airtable link once created)* · Live Prototype — Trip Assistant

---

The Problem

MakeMyTrip's recommendation engine kept surfacing the same popular, expensive, heavily-reviewed cities (Goa, Manali) regardless of what travellers actually typed — even an explicit offbeat request like "Spiti Valley" returned Manali resorts. This eroded trust and pushed travellers to plan manually via blogs/Reddit instead of using the app.

Research & Discovery

Built two proto-personas from behavioral evidence, not assumptions:

- Anjali Mehta ("The Offbeat Explorer") — budget-conscious, trust-eroded by biased recommendations. Her latent need (uncovered via empathy mapping, distinct from her stated request) was *transparency* — she needed to see *why* a place was recommended, not just a different answer.
- Rajeev Nair ("The Time-Starved Family Planner") — needed a *guaranteed hard-budget cap*, not just a faster tool.

Ran a competitive teardown (Goibibo as direct competitor, generic LLM chatbots as indirect) and scored four candidate features with RICE. Selected: a Conversational Trip Assistant with clarifying follow-up questions and explainable recommendations (RICE score 3,200 — highest of four candidates).

What I Built

- PRD-lite with INVEST-checked user stories, a North Star metric (% of users who save/book >=1 recommended item within 24h) and a guardrail metric (budget-violation rate, capped below 5%).
- Working no-code prototype (built via Google AI Studio / Lovable): email/password auth, persisted preferences + generated itineraries per user, itinerary generation constrained to a fixed catalogue (no hallucinated hotels), and a one-line "why this itinerary" explanation shown to the user.
- Automation pipeline (trigger -> format -> LLM call -> human-review checkpoint -> Airtable log): every AI-generated explanation is logged with a `Review Status` and held back from the user until a human reviewer approves it — a deliberate non-automated checkpoint because showing an unreviewed AI explanation directly to a customer is a hard-to-reverse action.
- Eval suite: 8-input golden dataset scored against a 5-criteria rubric (accuracy, clarity, completeness, tone, formatting) plus 5 binary assertion checks (budget respected, dates in window, no invented catalogue items, trip-type appropriate, explanation present).

Responsible AI

Designed four targeted bias mitigations tied directly to the prototype's ranking logic: - Popularity bias -> rotation cap on top-5 catalogued destinations in any 7-day window. - Price bias -> hard budget filter applied *before* the LLM sees candidates, not just de-prioritized after. - Geographical bias -> prompt requires >=2 geographic clusters when "offbeat" is selected. - Review bias -> blended ranking score (relevance weighted 2x, review-count weighted only 0.5x).

Plus standard mitigations for prompt injection (user text treated as data, never as instructions), hallucination (catalogue-only generation), and PII handling (scoped per-user, hashed auth).

Pricing, GTM & Growth

- Good-Better-Best pricing (Free / Explorer ₹149 / Concierge ₹399) with a decoy tier and anchoring against the ₹1,500+ cost of a human travel agent.
- Funnel diagnostic: identified the real bottleneck wasn't top-of-funnel awareness (30,000 users lost at banner->open) but the 70% drop-off between itinerary-viewed and booking/save — the trust-driven step this whole feature was designed to fix.
- AI Product Loop: shared/exported itineraries bring in new users via WhatsApp, who then generate and share their own.
- Bottom-up revenue model: ~₹1.08 crore projected Year-1 revenue at 100K MAU and 4% paid conversion.

Outcome

A fully designed, prototyped, and evaluated AI feature — from persona research through a working deployed app, an automated human-reviewed explanation pipeline, bias-mitigation design, and a GTM/pricing plan — grounded throughout in the specific trust and budget-compliance failures uncovered in user research, not generic AI-feature best practices.

Full case study (Parts 1–4, detailed PRD, Airtable base, eval scores, and automation transcript) available on request.